

# Review of Weight Uncertainty in Neural Network

Some useful links:

1. [Paper can be found here](#)
2. [Implementation by Pytorch on Github](#)

## Bayes by Backprop (BBP) Framework

A probabilistic model:  $P(y|x, w)$ : given an input  $x \in \mathbb{R}^p$ ,  $y \in \mathcal{Y}$ , using the set of parameters or weights  $w$ .

$$P(w|x, y) = \frac{P(y|x, w)p(w)}{P(y)}$$

## Loss Function

The weights can be learnt by MLE given a set of training samples  $\mathcal{D} = \{x_i, y_i\}_i$

$$\begin{aligned} w^{MLE} &= \arg \max_w \log P(\mathcal{D}|w) \\ &= \arg \max_w \log \sum_i^n P(y_i|x_i, w) \end{aligned}$$

Regularization can be done by add a prior on the weights  $w$  and finding the MAP, i.e.,

$$\begin{aligned} w^{MAP} &= \arg \max_w \log P(w|\mathcal{D}) \\ &= \arg \max_w \log P(\mathcal{D}|w)p(w) \end{aligned}$$

**Inference** is intractable because it needs to consider each configuration of  $w$ .

$$P(\hat{y}|\hat{x}) = \mathbb{E}_{p(w|\mathcal{D})}[P(\hat{y}|\hat{x}, w)]$$

# Minimization of KL Divergence (ELBO)

$$\begin{aligned}
 \theta^* &= \arg \min_{\theta} KL[q(w|\theta)|P(w|\mathcal{D})] \\
 &= \arg \min_{\theta} \int q(w|\theta) \log \frac{q(w|\theta)}{P(w|\mathcal{D})} dw \\
 &= \arg \min_{\theta} \int q(w|\theta) \log \frac{q(w|\theta)}{P(w)P(\mathcal{D}|w)} dw \\
 &= \arg \min_{\theta} \underbrace{KL[q(w|\theta)|P(w)]}_{\text{complexity cost}} - \underbrace{\mathbb{E}_{q(w|\mathcal{D})}[\log P(\mathcal{D}|w)]}_{\text{likelihood cost}}
 \end{aligned}$$

We denote it as:

$$\mathcal{F}(\mathcal{D}, \theta) = KL[q(w|\theta)|P(w)] - \mathbb{E}_{q(w|\theta)}[\log P(\mathcal{D}|w)] \quad (1)$$

$$\approx \frac{1}{n} \sum_{i=1}^n \log q(w^i|\theta) - \log P(w^i) - \log P(\mathcal{D}|w^i) \quad (2)$$

where  $w^i$  is a sample from variational posterior  $q(w|\theta)$ . Note that the parameters require gradient is  $\theta$  instead of  $w$  in MLE or MAP. Note that in the original paper there is no  $\frac{1}{n}$ . I think it is probably a typo.

## Unbiased Monte Carlo Gradients

Proposition 1. Let  $\epsilon$  be a random variable having probability density given by  $q(\epsilon)$  and let  $w = t(\theta, \epsilon)$  where  $t$  is a deterministic function. Suppose further that the marginal probability density of  $w$ ,  $q(w|\theta)$ , is such that  $q(\epsilon)d\epsilon = q(w|\theta)dw$ . Then for a function  $f$  with derivative in  $w$ :

$$\frac{\partial}{\partial \theta} \mathbb{E}_{q(w|\theta)}[f(w, \theta)] = \mathbb{E}_{q(\epsilon)} \left[ \frac{\partial f(w, \theta)}{\partial w} \frac{\partial w}{\partial \theta} + \frac{\partial f(w, \theta)}{\partial \theta} \right]$$

Our objective function in Eq. (1) can be written as

$$\begin{aligned}
 \mathcal{F}(\mathcal{D}, \theta) &= \mathbb{E}_{q(w|\theta)}[\log q(w|\theta) - \log P(w) - \log P(\mathcal{D}|w)] \\
 &= \mathbb{E}_{q(w|\theta)}[f(w, \theta)]
 \end{aligned}$$

We need the gradient

$$\begin{aligned}
\nabla_{\theta} \mathbb{E}_{q(w|\theta)}[f(w, \theta)] &= \nabla_{\theta} \int f(w, \theta) q(w|\theta) dw \\
&= \nabla_{\theta} \int f(w, \theta) q(\epsilon) d\epsilon \\
&= \mathbb{E}_{q(\epsilon)}[\nabla_{\theta} f(w, \theta)]
\end{aligned}$$

Leibniz integral rule can also be used here.

## Why re-parameterization trick

Some useful links:

1. [Introduction about MC gradient](#)
2. [Why re-parameterization trick](#)

Normally, we use Monte Carlo (MC) gradient estimator to approximate the gradient. It is used to solve the problem  $\nabla_{\theta} \mathbb{E}_{q(w)}[f(w, \theta)] = \mathbb{E}_{q(w)}[\nabla_{\theta} f(w, \theta)]$ . **Note that there is no parameter  $\theta$  in the expectation distribution  $q(w)$ , which is the difference between this case and the case we are facing.**

If we still use MC gradient estimator on  $\nabla_{\theta} \mathbb{E}_{q(w|\theta)}[f(w, \theta)]$ :

$$\begin{aligned}
\nabla_{\theta} \mathbb{E}_{q(w|\theta)}[f(w, \theta)] &= \int \nabla_{\theta} [q(w|\theta) f(w, \theta)] dw \\
&= \int \nabla_{\theta} [q(w|\theta)] f(w, \theta) dw + \int q(w|\theta) \nabla_{\theta} [f(w, \theta)] dw \\
&= \int q(w|\theta) \nabla_{\theta} [\log q(w|\theta)] f(w, \theta) dw + \mathbb{E}_{q(w|\theta)}[\nabla_{\theta} [f(w, \theta)]] \\
&= \mathbb{E}_{q(w|\theta)} \left[ \nabla_{\theta} [\log q(w|\theta)] f(w, \theta) \right] + \mathbb{E}_{q(w|\theta)}[\nabla_{\theta} [f(w, \theta)]]
\end{aligned}$$

The problem is that the distribution in the first term is coupled with both  $w$  and  $\theta$ . The value of the first term also depends on  $\theta$  but  $\theta$  is what we are tuning. We need to detach  $\theta$  from the distribution. Then, we can apply the MC gradient estimator.

**Why  $q(\epsilon) d\epsilon = q(w|\theta) dw$ ?**

For deterministic mapping  $w = t(\epsilon, \theta)$ ,  $q(\epsilon) d\epsilon = q(w|\theta) dw$  holds.

$$q(w|\theta) \frac{dw}{d\epsilon} = q(\epsilon)$$

$$q(w|\theta) = Kq(\epsilon)$$

$$G(\epsilon) = Kq(\epsilon)$$

## Leibniz integral rule

Leibniz integral rule:

$$\frac{d}{dx} \left( \int_{a(x)}^{b(x)} f(x, t) dt \right) = f(x, b(x)) \cdot \frac{d}{dx} b(x) - f(x, a(x)) \cdot \frac{d}{dx} a(x) + \int_{a(x)}^{b(x)} \frac{\partial}{\partial x} f(x, t) dt$$

## Mini-batches

For each epoch of optimization, the training set is equally and randomly split into  $M$  batches  $\mathcal{D}_1, \dots, \mathcal{D}_M$ . The loss can be rewritten as

$$\mathcal{F}_i(\mathcal{D}, \theta) = \frac{1}{M} KL[q(w|\theta)|P(w)] - \mathbb{E}_{q(w|\theta)}[\log P(\mathcal{D}|w)]$$

I'll update more details later.